Title: Building a Human-Friendly FAQ Bot from Jupiter Help Centre

Context:
Jupiter's app has a Help Centre filled with FAQs that explain topics such as payments, cards, KYC, rewards, and more. Users, however, often prefer quick and conversational answers instead of navigating through static pages. The objective of this project is to scrape these FAQs and build an intelligent bot that can answer user questions in a simple, conversational, and accurate manner.

Problem Breakdown:

1. **Scrape the FAQs:**
    - Crawl the FAQ section of the Jupiter app or help website.
    - Structure the scraped data into question–answer pairs, organized by category.
2. **Preprocess and Clean:**
    - Normalize and deduplicate similar questions.
    - Clean any HTML or formatting noise.
    - Categorize the content into topics such as KYC, rewards, payments, limits, etc.
3. **Build the FAQ Bot:**
    - Utilize large language models (LLMs) such as OpenAI, Mistral, or LLaMA 3 or any other LLM of your choice to:
        - Rephrase answers into friendly and natural language.
        - Handle semantically similar queries.
        - Respond confidently or gracefully decline if unsure.
    - Optionally integrate embedding-based semantic search using tools such as FAISS or Chroma or equivalent.
4. **Demonstrate the Solution:**
    - Develop a simple user interface or notebook to simulate user interactions.

    - The bot should retrieve the best match or generate an accurate, context-aware response.

Deliverables:

- A clean, well-documented notebook or application.
- Evaluation of semantic similarity and relevance of answers.
- Clear documentation of methodology and architecture.

Bonus Objectives:

- Implement multilingual support (e.g., Hindi or Hinglish).
- Add suggestion capabilities for related queries based on user behavior.
- Compare retrieval-based versus LLM-based approaches on accuracy and latency.